Introduction to Artificial Intelligence

Kalev Kask

ICS 271

Fall 2014

http://www.ics.uci.edu/~kkask/Fall-2014 CS271/

Course requirements

Assignments:

- There will be weekly homework assignments, a project, a final.

Course-Grade:

- Homework will account for 20% of the grade, project 30%, final 50% of the grade.

. **Discussion:**

- Optional. TBD.

Course overview

- Introduction and Agents (chapters 1,2)
- Search (chapters 3,4,5,6)
- Logic (chapters 7,8,9)
- Planning (chapters 10,11)
- Uncertainty and probability (chapters 13-14)

Plan of the course

Part I Artificial Intelligence

>   1 Introduction
>   2 Intelligent Agents

Part II Problem Solving

>   3 Solving Problems by Searching
>   4 Beyond Classical Search
>   5 Adversarial Search
>   6 Constraint Satisfaction Problems

Part III Knowledge and Reasoning

>   7 Logical Agents
>   8 First-Order Logic
>   9 Inference in First-Order Logic
>   10 Classical Planning
>   11 Planning and Acting in the Real World

Part IV Uncertainty

>   13 Uncertainly
>   14 Probabilistic Reasoning

Resources on the internet

Resources on the Internet

- [AI on the Web:](#) A very comprehensive list of Web resources about AI from the Russell and Norvig textbook.

Essays and Papers

- [What is AI](#), John McCarthy
- Computing Machinery and Intelligence, A.M. Turing
- [Rethinking Artificial Intelligence](#), Patrick H.Winston

- [AI Topics: http://aitopics.net/index.php](#)

Today's class

- What is Artificial Intelligence?
- A brief History
- Intelligent agents
- State of the art

Today's class

- What is Artificial Intelligence?

What is Artificial Intelligence
([John McCarthy](...) , Basic Questions

- **What is artificial intelligence?**
- It is the science and engineering of making intelligent machines, especially intelligent computer programs. It is related to the similar task of using computers to understand human intelligence, but AI does not have to confine itself to methods that are biologically observable.

- **Yes, but what is intelligence?**
- Intelligence is the computational part of the ability to achieve goals in the world. Varying kinds and degrees of intelligence occur in people, many animals and some machines.

- **Isn't there a solid definition of intelligence that doesn't depend on relating it to human intelligence?**
- Not yet. The problem is that we cannot yet characterize in general what kinds of computational procedures we want to call intelligent. We understand some of the mechanisms of intelligence and not others.

- **More in: http://www-formal.stanford.edu/jmc/whatisai/node1.html**

What is Artificial Intelligence?

- Human-like vs rational-like
- Thought processes vs behavior
- How to simulate humans intellect and behavior by a machine.
  - Mathematical problems (puzzles, games, theorems)
  - Common-sense reasoning
  - Expert knowledge: lawyers, medicine, diagnosis
  - Social behavior
- **Things we would call "intelligent" if done by a human.**

What is AI?

Views of AI fall into four categories:

Thinking humanly	Thinking rationally
Acting humanly	Acting rationally

The textbook advocates "acting rationally"

How to simulate humans intellect and behavior by a machine.
Mathematical problems (puzzles, games, theorems)
Common-sense reasoning
Expert knowledge: lawyers, medicine, diagnosis
Social behavior

The Turing Test
(Can Machine think? A. M. Turing, 1950)

http://aitopics.net/index.php

http://amturing.acm.org/acm_tcc_webcasts.cfm

THE INTERROGATOR

Figure 1.1 The Turing test.

- Requires:
  - Natural language
  - Knowledge representation
  - Automated reasoning
  - Machine learning
  - (vision, robotics) for full test

Acting/Thinking Humanly/Rationally

- Turing test (1950)
- Requires:
  - Natural language
  - Knowledge representation
  - automated reasoning
  - machine learning
  - (vision, robotics.) for full test
- Methods for Thinking Humanly:
  - Introspection, the general problem solver (Newell and Simon 1961)
  - Cognitive sciences
- Thinking rationally:
  - Logic
  - Problems: how to represent and reason in a domain
- Acting rationally:
  - Agents: Perceive and act

What is Artificial Intelligence

- Thought processes
  - "The exciting new effort to make computers **think** .. Machines with minds, in the full and literal sense" (Haugeland, 1985)

- Behavior
  - "The study of how to make computers **do things** at which, at the moment, people are better." (Rich, and Knight, 1991)

- Activities
  - The automation of activities that we associate with human thinking, activities such as decision-making, problem solving, learning… (Bellman)

271-fall 2014

The foundation of AI

Philosophy, Mathematics, Economics, Neuroscience, Psychology, Computer Engineering,

Features of intelligent system

- Deduction, reasoning, problem solving
- Knowledge representation
- Planning
- Learning
- Natural language processing
- Perception
- Motion and manipulation

Tools

- Search and optimization
- Logic
- Probabilistic reasoning
- Neural networks

Today's class

- What is Artificial Intelligence?
- A brief history
- Intelligent agents
- State of the art

Histroy of AI

- McCulloch and Pitts (1943)
  - Neural networks that learn
- Minsky and Edmonds (1951)
  - Built a neural net computer
- Darmouth conference (1956):
  - McCarthy, Minsky, Newell, Simon met,
  - Logic theorist (LT)- Of Newell and Simon proves a theorem in Principia Mathematica-Russel.
  - The name "Artficial Intelligence" was coined.
- 1952-1969 (early enthusiasm, great expectations)
  - GPS- Newell and Simon
  - Geometry theorem prover - Gelernter (1959)
  - Samuel Checkers that learns (1952)
  - McCarthy - Lisp (1958), Advice Taker, Robinson's resolution
  - Microworlds: Integration, block-worlds.
  - 1962- the perceptron convergence (Rosenblatt)

The Birthplace of "Artificial Intelligence", 1956

- **Darmouth workshop, 1956:** historical meeting of the precieved founders of AI met: John McCarthy, Marvin Minsky, Alan Newell, and Herbert Simon.

- **A Proposal for the Dartmouth Summer Research Project on Artificial Intelligence. J. McCarthy, M. L. Minsky, N. Rochester, and C.E. Shannon. August 31, 1955. "We propose that a 2 month, 10 man study of artificial intelligence be carried out during the summer of 1956 at Dartmouth College in Hanover, New Hampshire. The study is to proceed on the basis of the conjecture that every aspect of learning or any other feature of intelligence can in principle be so precisely described that a machine can be made to simulate it." *And this marks the debut of the term* "artificial intelligence."**

- **50 anniversery of Darmouth workshop**
- List of AI-topics

More AI examples

Common sense reasoning (1980-1990)
- Tweety
- Yale Shooting problem

Update vs revise knowledge

 The OR gate example: A or B → C
- Observe C=0, vs Do C=0

Chaining theories of actions

      Looks-like(P) → is(P)

      Make-looks-like(P) → Looks-like(P)

      Makes-looks-like(P) ---is(P) ???

Garage-door example: garage door not included.
- Planning benchmarks
- 8-puzzle, 8-queen, block world, grid-space world
- Cambridge parking example

Smoked fish example… what is this?

History, continued

- 1966-1974 a dose of reality
  - Problems with computation
- 1969-1979 Knowledge-based systems
  - Weak vs. strong methods
  - Expert systems:
    - Dendral:Inferring molecular structures(Buchanan et. Al. 1969)
    - Mycin: diagnosing blood infections (Shortliffe et. Al, certainty factors)
    - Prospector: recommending exploratory drilling (Duda).
  - Roger Shank: no syntax only semantics
- 1980-1988: AI becomes an industry
  - R1: Mcdermott, 1982, order configurations of computer systems
  - 1981: Fifth generation
- 1986-present: return to neural networks
- 1987-present :
  - **AI becomes a science**: HMMs, planning, belief network
- 1995-present: The emergence of intelligent agents
  - Ai agents (SOAR, Newell, Laired, 1987) on the internet, technology in web-based **applications** , recommender systems. Some researchers (Nilsson, McCarthy, Minsky, Winston) express discontent with the progress of the field. AI should return to human-level AI (they say).
- 2001-present: The availability of data;
  - The knowledge bottleneck may be solved for many applications: learn the information rather than hand code it .

State of the art

- **Game Playing:** Deep Blue defeated the reigning world chess champion Garry Kasparov in 1997
- **Robotics vehicles**:
  - 2005 Standford robot won DARPA Grand Challenge, driving autonomously 131 miles along unrehearsed desert trail
  - Staneley (Thrun 2006). No hands across America (driving autonomously 98% of the time from Pittsburgh to San Diego
  - 2007 CMU team won DARPA Urban Challenge driving autonomously 55 miles in a city while adhering to traffic laws and hazards
- **Autonomous planning and scheduling**:
  - During the 1991 Gulf War, US forces deployed an AI logistics planning and scheduling program that involved up to 50,000 vehicles, cargo, and people
  - NASA's on-board autonomous planning program controlled the scheduling of operations for a spacecraft
- **Speech recognition**
- DARPA grand challenge 2003-2005, Robocup
- **Machine translation** (From English to arabic, 2007)
- **Natural language processing**: Watson won Jeopardy (Natural language processing), IBM 2011. iPhone Siri.

Robotic links

- Deep Blue: http://en.wikipedia.org/wiki/Deep_Blue_(chess_computer)
- **Robocup Video**
  - Soccer Robocupf

- Darpa Challenge

  - Darpa's-challenge-video

- Watson
- http://www.youtube.com/watch?v=seNkjYyG3gI

Today's class

- What is Artificial Intelligence?
- A brief History
- Intelligent agents
- State of the art

271-fall 2014

Agents

- An agent is anything that can be viewed as perceiving its environment through sensors and acting upon that environment through actuators

- Human agent: eyes, ears, and other organs for sensors; hands, legs, mouth, and other body parts for actuators

- Robotic agent: cameras and infrared range finders for sensors; various motors for actuators

Agents

- Agents and environments
- Rationality
- PEAS (Performance measure, Environment, Actuators, Sensors)
- Environment types
- Agent types

Agents and environments

- The agent function maps from percept histories to actions:

$$[f: \mathcal{P}^{\star} \rightarrow \mathcal{A}]$$

- The agent program runs on the physical architecture to produce f

- agent = architecture + program

What's involved in Intelligence?

- **Ability to interact with the real world**
  - to perceive, understand, and act
  - e.g., speech recognition and understanding and synthesis
  - e.g., image understanding
  - e.g., ability to take actions, have an effect

- **Knowledge Representation, Reasoning and Planning**
  - modeling the external world, given input
  - solving new problems, planning and making decisions
  - ability to deal with unexpected problems, uncertainties

- **Learning and Adaptation**
  - we are continuously learning and adapting
  - our internal models are always being "updated"
    - e.g. a baby learning to categorize and recognize animals

Implementing agents

- **Table look-ups**
- **Autonomy**
  - All actions are completely specified
  - no need in sensing, no autonomy
  - example: Monkey and the banana
- **Structure of an agent**
  - agent = architecture + program
  - Agent types
    - medical diagnosis
    - Satellite image analysis system
    - part-picking robot
    - Interactive English tutor
    - cooking agent
    - taxi driver
    - **Graduate student**

Vacuum-cleaner world

Percepts: location and contents, e.g., $[A, Dirty]$

Actions: $Left$, $Right$, $Suck$, $NoOp$

A vacuum-cleaner agent

Percept sequence	Action
$[A, Clean]$	$Right$
$[A, Dirty]$	$Suck$
$[B, Clean]$	$Left$
$[B, Dirty]$	$Suck$
$[A, Clean], [A, Clean]$	$Right$
$[A, Clean], [A, Dirty]$	$Suck$
⋮	⋮

function REFLEX-VACUUM-AGENT($[location, status]$) **returns** an action

    if $status = Dirty$ **then return** $Suck$
    else if $location = A$ **then return** $Right$
    else if $location = B$ **then return** $Left$

What is the right function?
Can it be implemented in a small agent program?

Rationality

Fixed performance measure evaluates the environment sequence
- one point per square cleaned up in time T?
- one point per clean square per time step, minus one per move?
- penalize for $> k$ dirty squares?

A rational agent chooses whichever action maximizes the expected value of the performance measure given the percept sequence to date

Rational $\neq$ omniscient
Rational $\neq$ clairvoyant
Rational $\neq$ successful

Rational $\Rightarrow$ exploration, learning, autonomy

Task Environment

- Before we design a rational agent, we must specify its **task environment**:

PEAS:

Performance measure

Environment

Actuators

Sensors

PEAS

- Example: Agent = taxi driver

  – Performance measure: Safe, fast, legal, comfortable trip, maximize profits

  – Environment: Roads, other traffic, pedestrians, customers

  – Actuators: Steering wheel, accelerator, brake, signal, horn

  – Sensors: Cameras, sonar, speedometer, GPS, odometer, engine sensors, keyboard

PEAS

- Example: Agent = Medical diagnosis system

  – Performance measure: Healthy patient, minimize costs, lawsuits

  – Environment: Patient, hospital, staff

  – Actuators: Screen display (questions, tests, diagnoses, treatments, referrals)

  – Sensors: Keyboard (entry of symptoms, findings, patient's answers)

PEAS

- Example: Agent = Me Part-picking robot

  – Performance measure: Percentage of parts in correct bins

  – Environment: Conveyor belt with parts, bins

  – Actuators: Jointed arm and hand

  – Sensors: Camera, joint angle sensors

Environment Types

- **Fully observable** (vs. **partially observable**): An agent's sensors give it access to the complete state of the environment at each point in time.

- **Deterministic** (vs. **stochastic**): The next state of the environment is completely determined by the current state and the action executed by the agent. (If the environment is deterministic except for the actions of other agents, then the environment is **strategic**)

- **Episodic** (vs. **sequential**): An agent's action is divided into atomic episodes. Decisions do not depend on previous decisions/actions.

Environment Types

- Static (vs. dynamic): The environment is unchanged while an agent is deliberating. (The environment is semidynamic if the environment itself does not change with the passage of time but the agent's performance score does)

- Discrete (vs. continuous): A limited number of distinct, clearly defined percepts and actions.

   How do we represent or abstract or model the world?

- Single agent (vs. multi-agent): An agent operating by itself in an environment. Does the other agent interfere with my performance measure?

Environment types

	Solitaire	Backgammon	Internet shopping	Taxi
Observable??				
Deterministic??				
Episodic??				
Static??				
Discrete??				
Single-agent??				

Environment types

	Solitaire	Backgammon	Internet shopping	Taxi
Observable??	Yes	Yes	No	No
Deterministic??				
Episodic??				
Static??				
Discrete??				
Single-agent??				

Environment types

	Solitaire	Backgammon	Internet shopping	Taxi
Observable??	Yes	Yes	No	No
Deterministic??	Yes	No	Partly	No
Episodic??				
Static??				
Discrete??				
Single-agent??				

271-fall 2014

Environment types

	Solitaire	Backgammon	Internet shopping	Taxi
Observable??	Yes	Yes	No	No
Deterministic??	Yes	No	Partly	No
Episodic??	No	No	No	No
Static??				
Discrete??				
Single-agent??				

Environment types

	Solitaire	Backgammon	Internet shopping	Taxi
Observable??	Yes	Yes	No	No
Deterministic??	Yes	No	Partly	No
Episodic??	No	No	No	No
Static??	Yes	Semi	Semi	No
Discrete??				
Single-agent??				

Environment types

	Solitaire	Backgammon	Internet shopping	Taxi
Observable??	Yes	Yes	No	No
Deterministic??	Yes	No	Partly	No
Episodic??	No	No	No	No
Static??	Yes	Semi	Semi	No
Discrete??	Yes	Yes	Yes	No
Single-agent??				

Environment types

Grad student

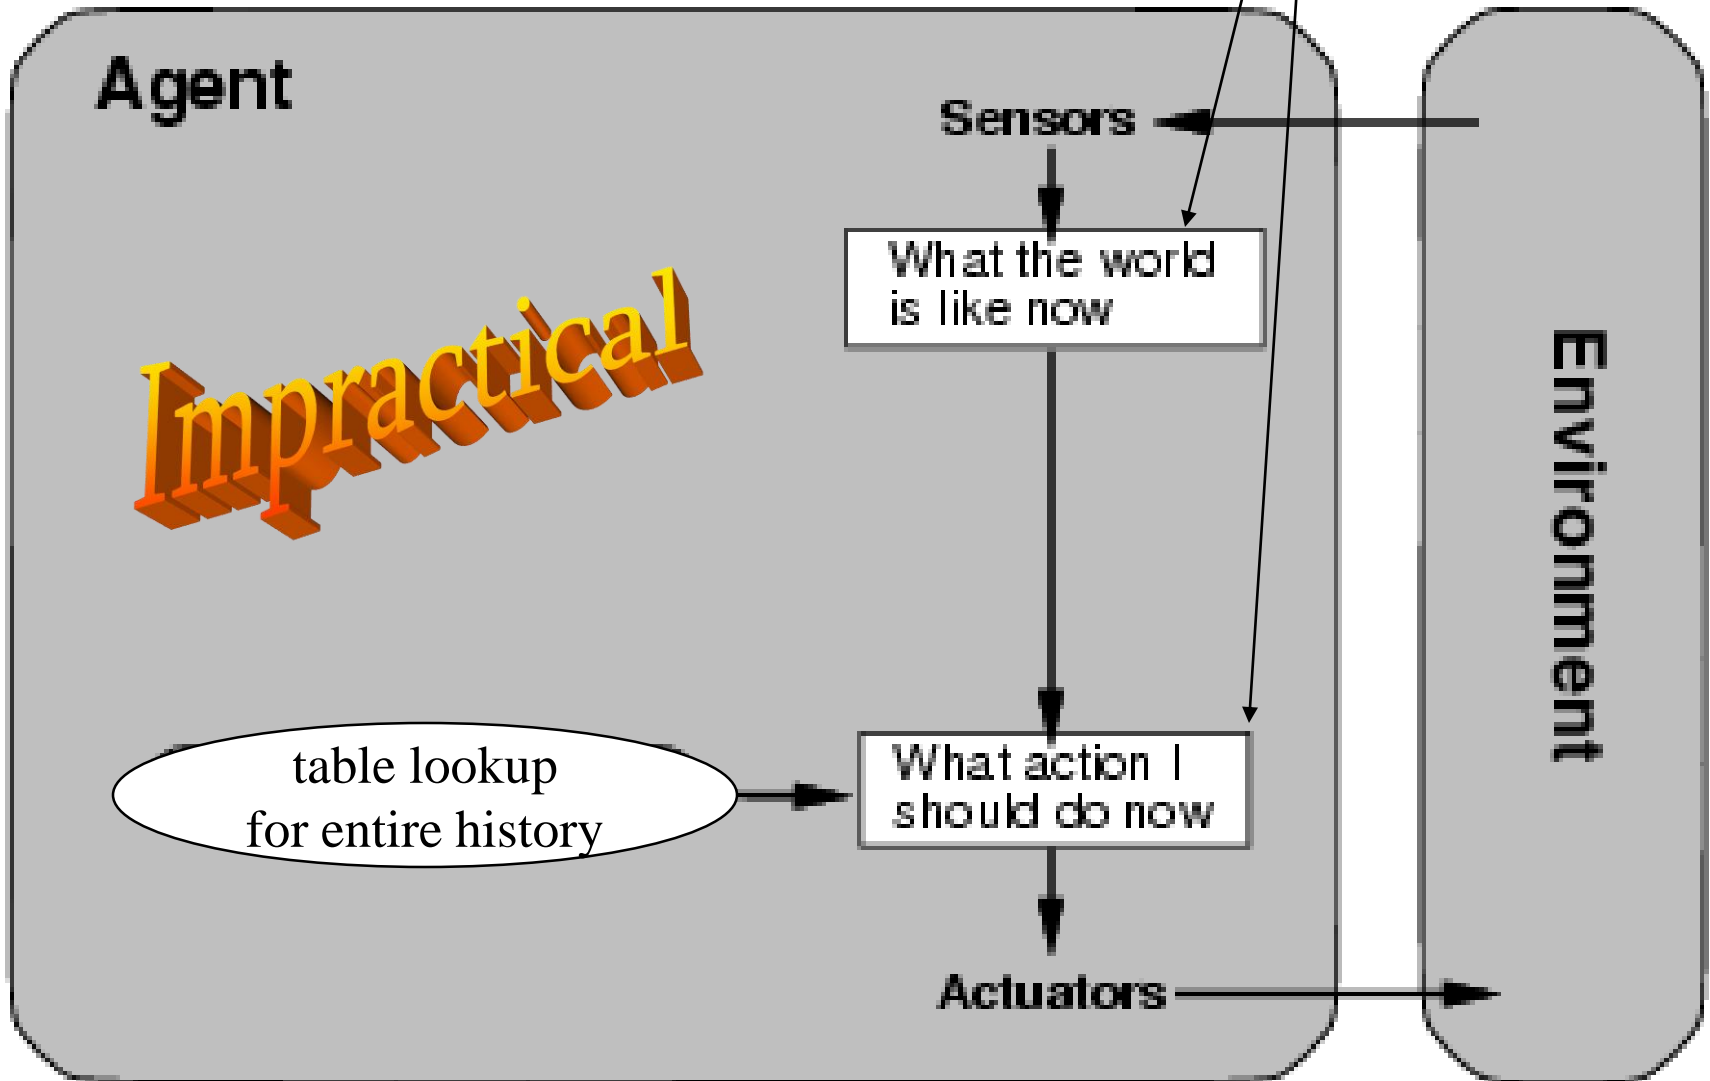|	Solitaire	Backgammon	Internet shopping	Taxi
Observable??	Yes	Yes	No	No
Deterministic??	Yes	No	Partly	No
Episodic??	No	No	No	No
Static??	Yes	Semi	Semi	No
Discrete??	Yes	Yes	Yes	No
Single-agent??	Yes	No	No (except auctions)	No

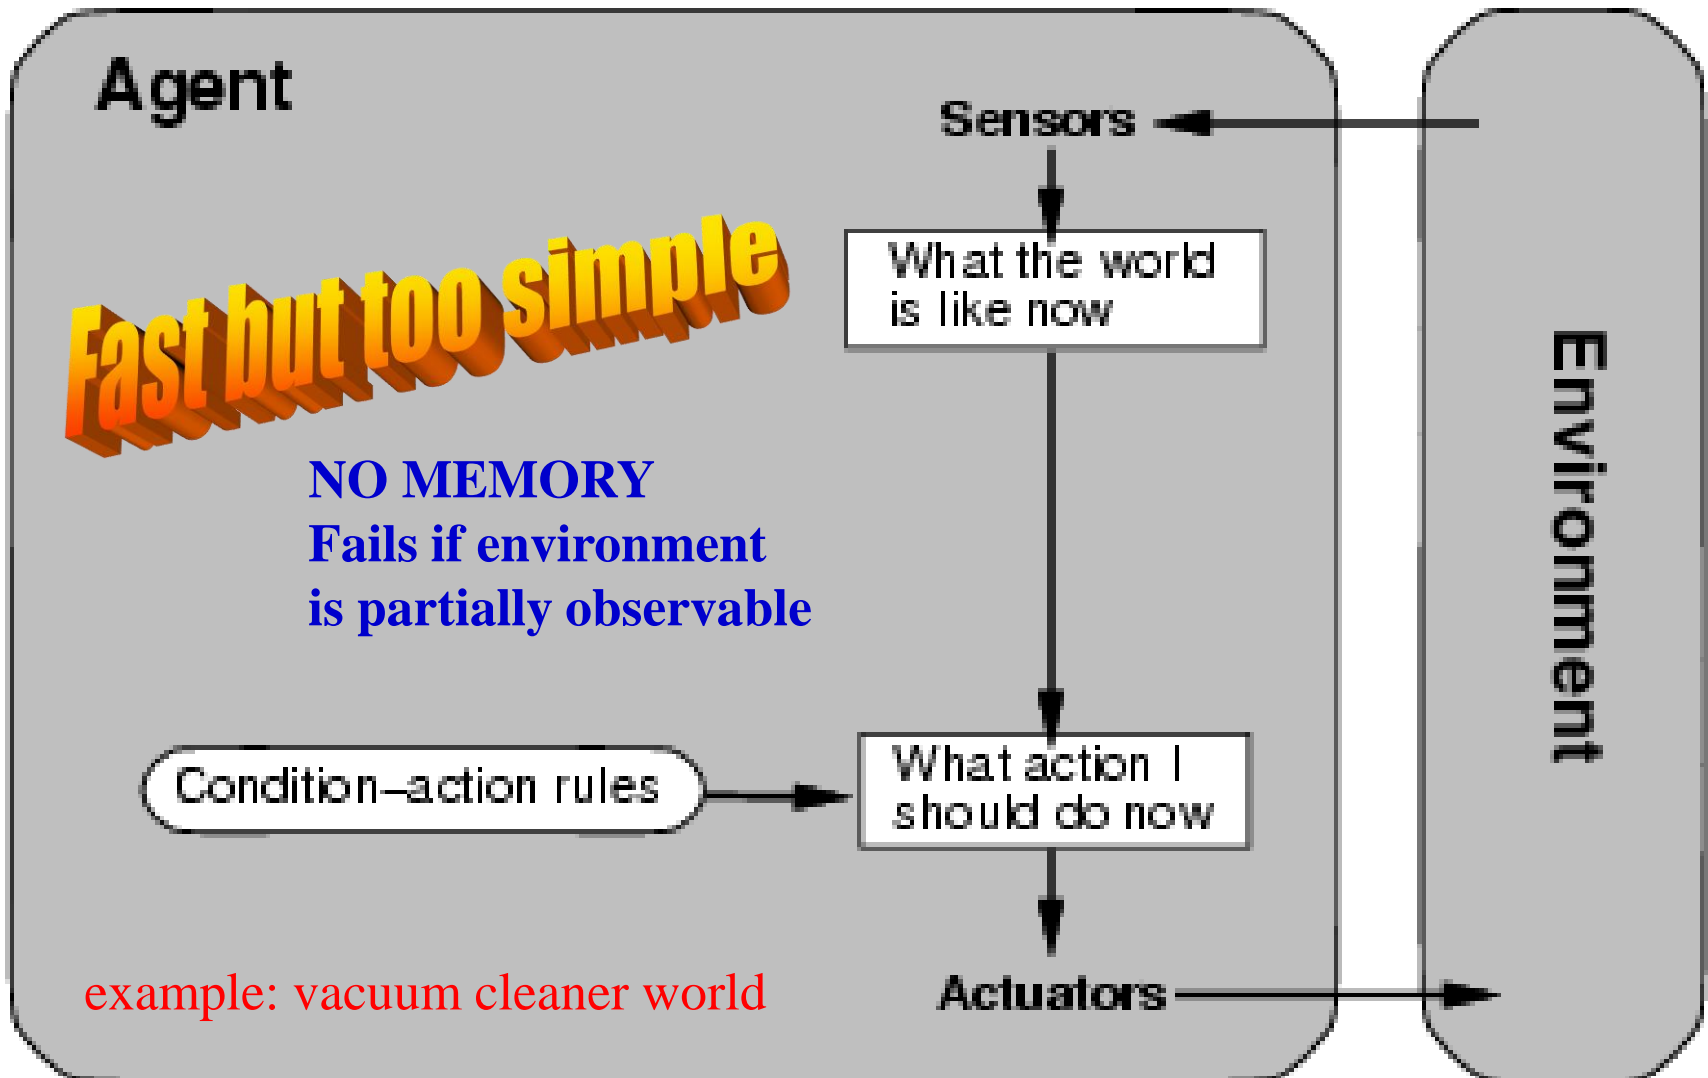The environment type largely determines the agent design

The real world is (of course) partially observable, stochastic, sequential, dynamic, continuous, multi-agent

271-fall 2014

Table Driven Agent.

Agent

Environment

Sensors

What the world is like now

Impractical

table lookup
for entire history

What action I
should do now

Actuators

Simple reflex agents

Agent

Fast but too simple

NO MEMORY
Fails if environment
is partially observable

Sensors

What the world is like now

Condition–action rules → What action I should do now
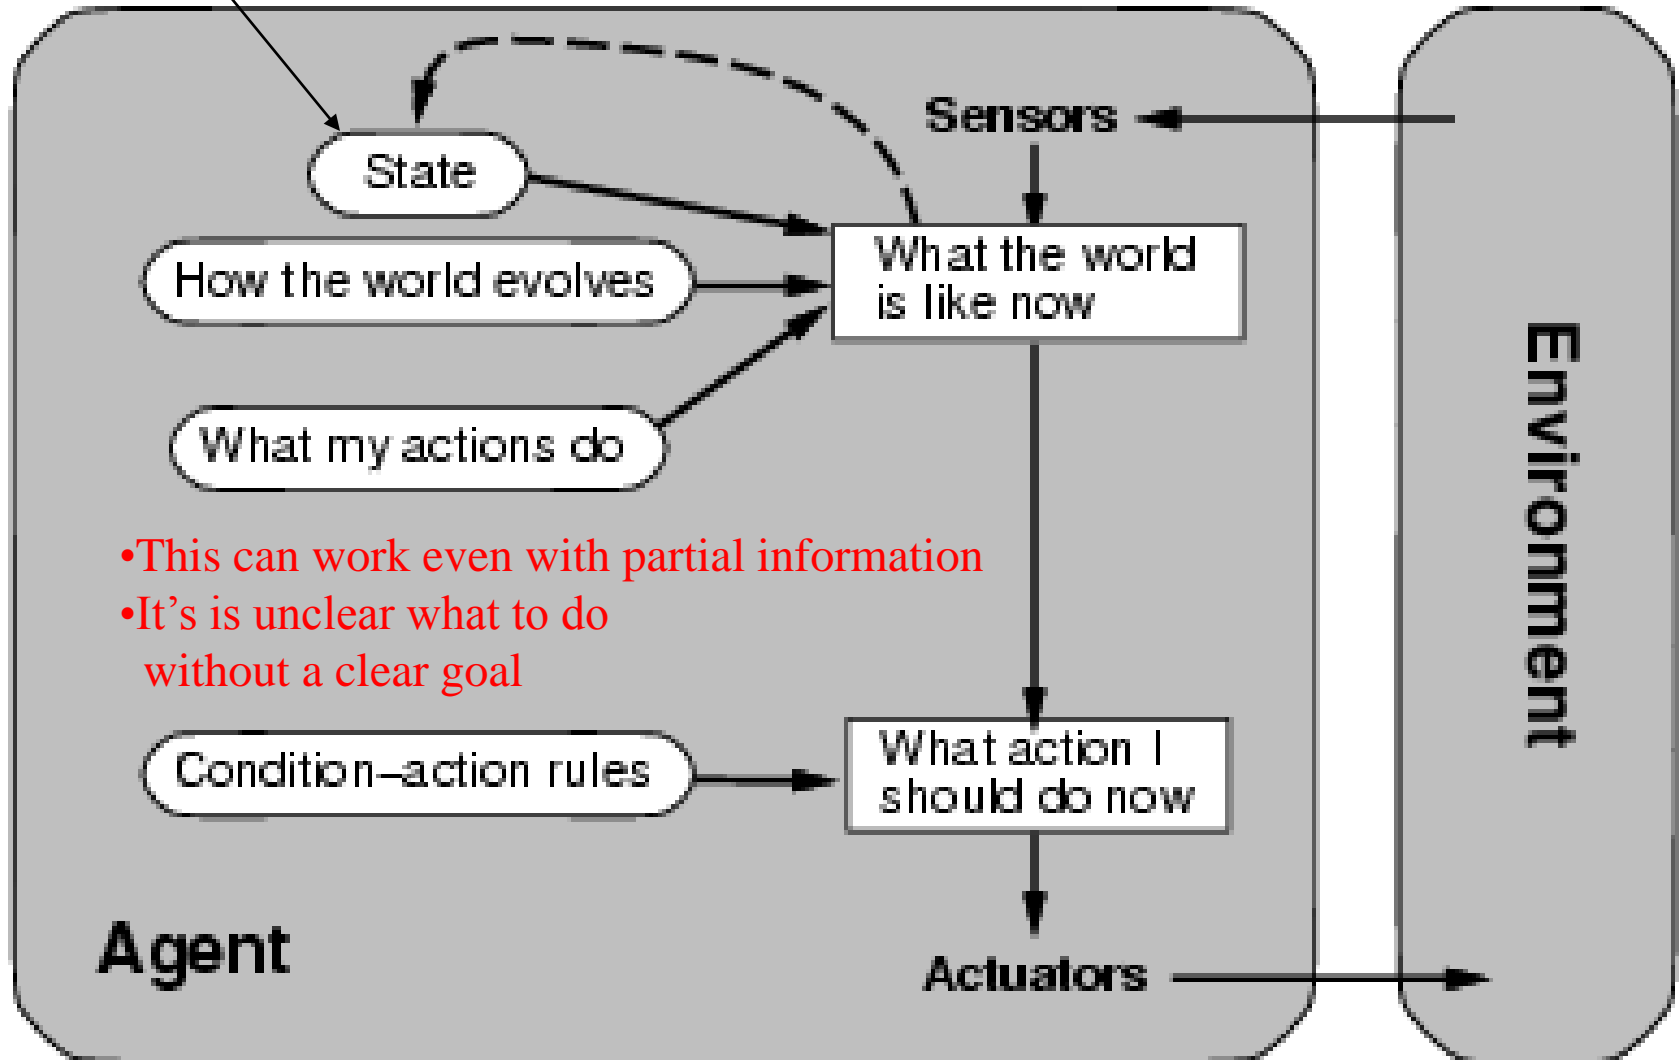
example: vacuum cleaner world

Actuators

Environment

Model-based reflex agents

description of
current world state

Model the state of the world by:
modeling how the world changes
how it's actions change the world

Sensors

State

How the world evolves

What the world
is like now

What my actions do

•This can work even with partial information
•It's is unclear what to do
 without a clear goal

Condition–action rules

What action I
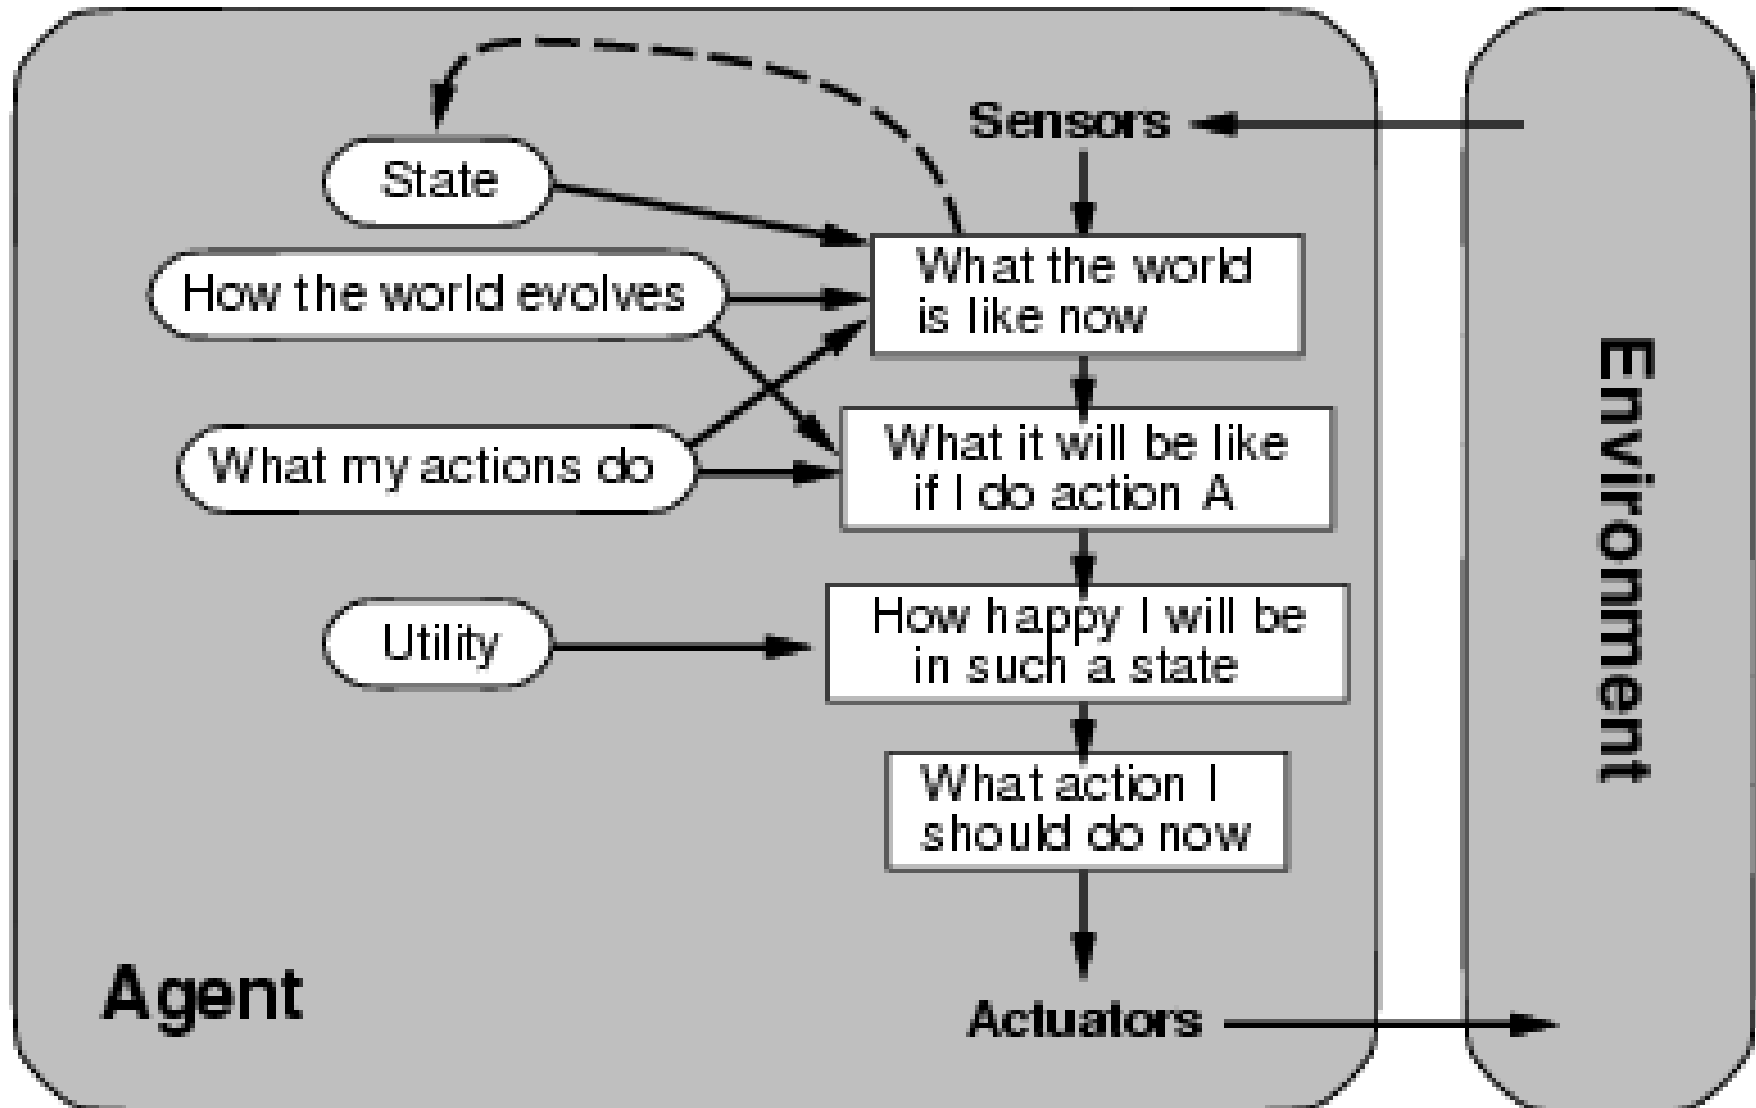should do now

Agent

Actuators

Environment

Goal-based agents

Goals provide reason to prefer one action over the other.
We need to predict the future: we need to plan & search

Utility-based agents

Some solutions to goal states are better than others.
Which one is best is given by a utility function.
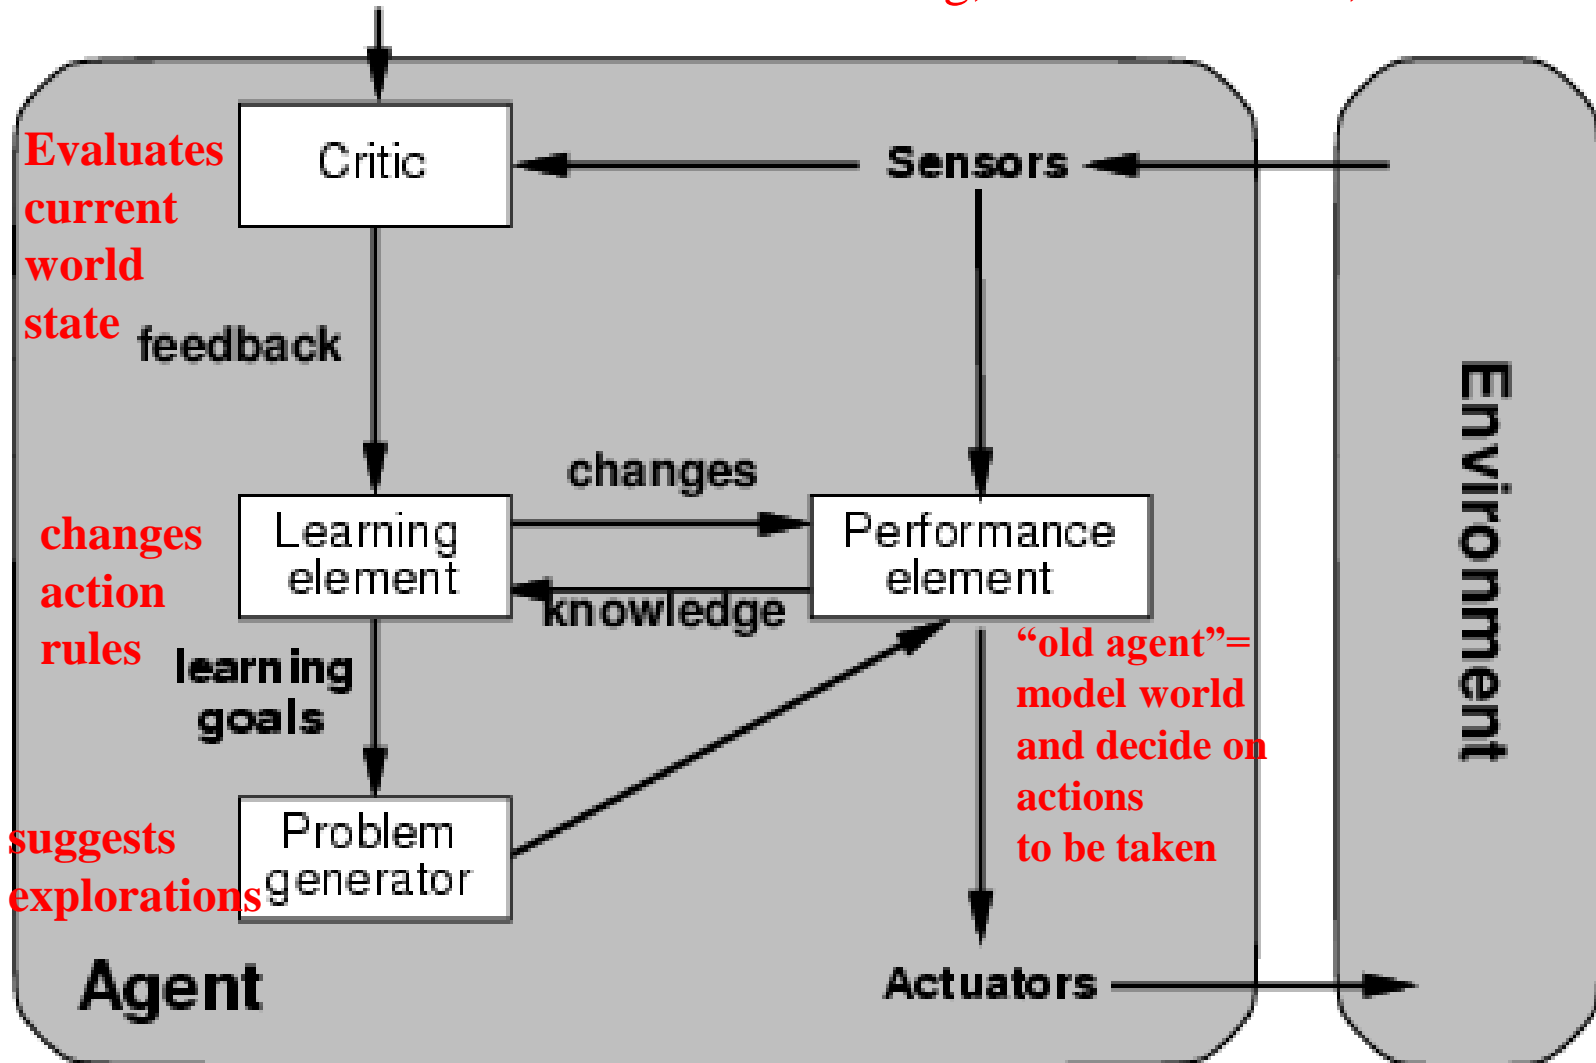Which combination of goals is preferred?

Learning agents

How does an agent improve over time?
By monitoring it's performance and suggesting
better modeling, new action rules, etc.

Summary

- **What is Artificial Intelligence**?
  - modeling humans thinking, acting, should think, should act.
- **History of AI**
- **Intelligent agents**
  - We want to build agents that act rationally

- **Real-World Applications of AI**
  - AI is alive and well in various "every day" applications
    - many products, systems, have AI components
- **Assigned Reading**
  - Chapters 1 and 2 in the text R&N